

Maschmeyer's Chosen Portfolio

You don't apply. You get chosen.

Early-signal deal sourcing: public builder signals → golden people → ranked ventures

One warehouse. One person_id. Every score carries its evidence.

SOURCES · WHAT WE SCRAPE

GitHub

Newly created trending repos,
core contributors, commit history

arXiv + OpenAlex

Papers, citations, affiliations;
Semantic Scholar, Papers-with-Code

Hack Nation

our own connector

People + projects JSON API,
teams, pitches, CVs

Zefix / SHAB

Swiss registry — schema ready

next

BRONZE · RAW

Immutable landing zone

One row per source object, kept
whole as a VARIANT payload.

Idempotent by construction

Parquet staged to a UC Volume,
then a single MERGE.
content_hash skips unchanged rows,
so re-runs are free.

Nothing crashes a run

Validation failures land in
bronze._rejects and the sweep
continues.

Erasure suppression is enforced
inside the MERGE — a re-scrape
cannot resurrect a deleted person.

SILVER · IDENTITY

person_source_record

one identity per source — facts hang here

↓ person_source_link

person_id

one per human, whatever the source

Every downstream table joins on it —
contributions, papers, roles, venture teams.

Reversible by design

Links are append-only, so unmerging
a mistake never touches a fact row.

How we decide

8 deterministic rules → Splink →
Claude settles the ambiguous band.

GOLD · SCORING

venture

a repo, a company or a hackathon
project — plus the people behind it

members joined on person_id

9 category scores

Experience, schools, network, prior
collaboration, problem, product,
market, traction, ideal-candidate fit.

Confidence, not just a number

Thin evidence lowers confidence and
flags the venture for more data.

Cited memo

Every bullet carries a source URL,
or is marked missing on purpose.

APP · THE LOOP

Lovable UI

Ranked list, memo, weight sliders,
ideal-candidate editor

Edge proxy · /v1

Holds the OAuth credential server-
side; the browser never sees a token

Reads gold.v_* views only, never base
tables. Precomputed, so pages load fast.

The loop that closes it

Shortlist → consent-based outreach →
AI interview fills the known gaps →
rescore → the memo rewrites itself.

Move a weight slider and the list
re-ranks in the browser — no job runs.

DATABRICKS

Unity Catalog · Delta · serverless SQL warehouse · ai_query · gte-large-en embeddings · UC Volumes

LOVABLE CLOUD

Supabase

How we score — and why an embedding on its own would be wrong

Everything below is computed inside the warehouse, over the same feature layer, and every weight is the fund's to change.

1 · The feature layer

Ten signals per person, deterministic
first, computed in SQL:

stars_weighted = log1p(Σ stars × share)
recency = 0.5 ^ (days / 90)
commits_12mo · graph centrality
school_tier · citations · founder flag

A missing signal stays NULL — never 0.

2 · Embeddings, deliberately limited

databricks-gte-large-en, 1024 dimensions,
normalized, compared as a dot product
in plain SQL. No vector database.

We use them for one thing only: does this
person work on that kind of problem?

An embedding cannot know that MIT outranks
KTH, or that 8,200 stars dwarf 82 — so
prestige and scale live in the feature layer.

3 · From person to team to rank

final = Σ w · s / Σ w_available
team = 0.6 · mean + 0.4 · max
conf = Σ w · coverage · diversity / Σ w

One strong founder should lift a team without
hiding a weak one, so we blend the average
with the best member.

A category with no evidence hands its weight to the rest.

4 · Where Claude does real work

Called from inside SQL, over the whole
table at once, to judge what a query
cannot: is this commit history any good,
is this product actually defensible,
are these two records the same person.

For the shortlist, Opus researches the
market on the open web and must quote
its sources — capped and cached.